

Capture-Recapture Example (Tutorial)

Leah South

Our running example is about the capture and recapture of the bird species called the European Dipper (*Cinclus cinclus*). Marzolin (1988) collected the data based on the capture and recapture of this species over six years.



Below are the R packages we'll be using in this document.

```
library(MASS) # multivariate normal
library(coda) # assessing convergence and sample quality
library(psych) # bivariate plots
```

The MCMC Algorithm

We'll be looking at MALA (asymptotically un-biased) and ULA (biased) approaches for inference on our example. The focus is on post-processing and bias rather than implementation of algorithms. You can find the relevant code for MALA and ULA implementations for reference.

```
source('MCMC_fn.R')
source('ULA_fn.R')
```

The Statistical Model

The parameters for the model are ϕ_i and p_k where $i = 1, \dots, 6$ and $k = 2, \dots, 7$. ϕ_i represents the probability of survival from year i to year $i + 1$ and p_k represents the probability of being captured in year k .

The likelihood for the model is given below, and based on data D_i for the number of birds released in year i and y_{ik} for the number of animals caught in year k out of the number released in year i . Here $d_i = D_i - \sum_{k=i+1}^7 y_{ik}$ is the number released in year i that are never caught. The corresponding probability of a bird being released in year i and never being caught is $\chi_i = 1 - \sum_{k=i+1}^7 \phi_i p_k \prod_{m=i+1}^{k-1} \phi_m (1 - p_m)$, which is a function of the model parameters. The likelihood is given by

$$f(y|\theta) \propto \prod_{i=1}^6 \chi_i^{d_i} \prod_{k=i+1}^7 \left[\phi_i p_k \prod_{m=i+1}^{k-1} \phi_m (1 - p_m) \right]^{y_{ik}},$$

where $\theta = (\phi, p)$, $\phi = (\phi_1, \dots, \phi_6)$, $p = (p_2, \dots, p_7)$ and $y = \{y_{ik} : i = 1, \dots, 6, k = 2, \dots, 7\}$. Due to parameter identifiability issues, the parameters ϕ_6 and p_7 are combined as $\phi_6 p_7$ leading to a total of eleven parameters.

The prior for each component of θ is set to be $\mathcal{U}(0, 1)$, and all components are independent *a priori*. For the RW proposal, the j -th parameter $\theta[j]$ is transformed using $\tilde{\theta}[j] = \log(\theta[j]/(1 - \theta[j]))$ for $j = 1, \dots, 11$. The implied prior density for $\tilde{\theta}[j]$ is then $e^{\tilde{\theta}[j]}/(1 + e^{\tilde{\theta}[j]})^2$, for $j = 1, \dots, 11$.

Read in the likelihood functions and tuning parameters for algorithms

```
load("recapture_ULA_bettertuning.RData")
load("recapture_MALA_tuning.RData")

# Names of the 11 variables
varNames <- paste0("theta", 1:11) ##### FIX-LATER
```

Multiple Chains

Getting the samples

Let's run 10 chains with a common starting point

```
initial <- c(0.35, -0.66, -1.74, 2.5, -0.67, -0.59, 2.38, 2.52, 1.2, 5.08,
1.3)
set.seed(2)
n_reps <- 10 # number of chains
its <- 500 # number of MCMC iterations
chains_ULA <- chains_MALA <- samples_ULA <- samples_MALA <- list()
for (i in 1:n_reps){

  # Running MALA
  single_chain_mala <- MALA_fn(d = 11, initial = initial, covmala = cov_rw, h = h_mala,
                              iters = its, der_loglike = der_loglike,
                              der_logprior = der_logprior,
                              options = options, varNames = varNames)

  chains_MALA[[i]] <- single_chain_mala
  samples_MALA[[i]] <- single_chain_mala$samples

  # Running ULA
  single_chain_ula <- ULA_fn(d = 11, initial = initial, cov_ULA = cov_ula, h = h_ula,
                             iters = its, der_loglike = der_loglike,
                             der_logprior = der_logprior,
                             options = options, varNames = varNames)

  chains_ULA[[i]] <- single_chain_ula
  samples_ULA[[i]] <- single_chain_ula$samples

}
ula_noburnin <- as.mcmc.list(samples_ULA)
mala_noburnin <- as.mcmc.list(samples_MALA)
save(ula_noburnin, mala_noburnin, chains_MALA, chains_ULA, file = "ten_chains.RData") # For future reuse
```

Compare the KSD

Sourcing in the KSD code and load the KSD package

```
source('KSD.R')
library(KSD)
```

Warning: package 'KSD' was built under R version 4.0.5

Evaluate the KSD on each of the chains for MALA and ULA.

```
KSD_MALA_gaussian <- KSD_ULA_gaussian <- rep(NA, n_reps)
KSD_MALA_imq <- KSD_ULA_imq <- rep(NA, n_reps)
for (i in 1:n_reps){
  samples <- as.matrix(chains_MALA[[i]]$samples)
  gradients <- chains_MALA[[i]]$der_loglike + chains_MALA[[i]]$der_logprior

  KSD_MALA_gaussian[i] <- KSD(samples, gradients)$ksd
  KSD_MALA_imq[i] <- imqKSD(samples, gradients)


  samples <- as.matrix(chains_ULA[[i]]$samples)
```

```

gradients <- chains_ULA[[i]]$der_loglike + chains_ULA[[i]]$der_logprior

KSD_ULA_gaussian[i] <- KSD(samples, gradients)$ksd
KSD_ULA_imq[i] <- imqKSD(samples, gradients)
}
save(KSD_MALA_gaussian,KSD_ULA_gaussian,KSD_MALA_imq,KSD_ULA_imq, file = "ksd_ten_chains.RData")
boxplot(KSD_MALA_gaussian,KSD_ULA_gaussian)


```



```

boxplot(KSD_MALA_imq,KSD_ULA_imq)

```



Estimating expectations with control variates

Now we'll estimate the posterior expectation of our parameters. The parameters are transformed using $\tilde{\theta}[j] = \log(\theta[j]/(1 - \theta[j]))$ for $j = 1, \dots, 11$. To transform back $\tilde{\theta}[j]$ we use $e^{\tilde{\theta}[j]}/(1 + e^{\tilde{\theta}[j]})^2$, for $j = 1, \dots, 11$.

```
library(ZVCV)
```

```
## Warning: package 'ZVCV' was built under R version 4.0.4
```

```
Vanilla_MALA <- ZV1_MALA <- CF_MALA <- SECF_MALA <- matrix(NaN,nrow=n_reps,ncol=d)
```

```
Vanilla_ULA <- ZV1_ULA <- CF_ULA <- SECF_ULA <- matrix(NaN,nrow=n_reps,ncol=d)
```

```
for (i in 1:n_reps){
```

```
  samples <- as.matrix(chains_MALA[[i]]$samples)
```

```
  gradients <- chains_MALA[[i]]$der_loglike + chains_MALA[[i]]$der_logprior
```

```
  integrand <- 1/(1+exp(-samples))
```

```
  # In order: vanilla estimate, zero-variance control variates
```

```
  # with a first order polynomial, control functionals and
```

```
  # semi-exact control functionals with a first order polynomial
```

```
  Vanilla_MALA[i,] <- colMeans(integrand)
```

```
  ZV1_MALA[i,] <- zvcv(integrand, samples, gradients,
```

```
                        options = list(polyorder = 1, regul_reg = FALSE))$expectation
```

```
  CF_MALA[i,] <- CF_crossval(integrand, samples, gradients, kernel_function = "RQ",
```

```
                           sigma_list = list(0.001,0.01,0.1,1,10), folds = 2)$expectation
```

```
  SECF_MALA[i,] <- SECF_crossval(integrand, samples, gradients, polyorder = 1, kernel_function = "RQ",
```


```
                                sigma_list = list(0.001,0.01,0.1,1,10), folds = 2)$expectation
```


```


# In order: vanilla estimate, zero-variance control variates
# with a first order polynomial, control functionals and
# semi-exact control functionals with a first order polynomial
samples <- as.matrix(chains_ULA[[i]]$samples)
gradients <- chains_ULA[[i]]$der_loglike + chains_ULA[[i]]$der_logprior
integrand <- 1/(1+exp(-samples))


Vanilla_ULA[i,] <- colMeans(integrand)
ZV1_ULA[i,] <- zvcv(integrand, samples, gradients,
                    options = list(polyorder = 1, regul_reg = FALSE))$expectation
CF_ULA[i,] <- CF_crossval(integrand, samples, gradients, kernel_function = "RQ",
                          sigma_list = list(0.001,0.01,0.1,1,10), folds = 2)$expectation
SECF_ULA[i,] <- SECF_crossval(integrand, samples, gradients, polyorder = 1, kernel_function = "RQ",
                              sigma_list = list(0.001,0.01,0.1,1,10), folds = 2)$expectation
}
load("Recapture_goldstandard.RData")
# Boxplots of the estimates
for (j in 1:11){
  boxplot(Vanilla_MALA[,j],ZV1_MALA[,j],CF_MALA[,j],SECF_MALA[,j],
          Vanilla_ULA[,j],ZV1_ULA[,j],CF_ULA[,j],SECF_ULA[,j])
  abline(h=gold_standard[j])
}


```





























Investigate tuning parameter for ULA

Try increasing and decreasing the parameter `h_ula`.

```
set.seed(2)
h_value <- 0.5
# Running ULA
single_chain_ula <- ULA_fn(d = 11, initial = initial, cov_ULA = cov_ula, h = h_value,
                           iters = its, der_loglike = der_loglike,
                           der_logprior = der_logprior,
                           options = options, varNames = varNames)
samples <- as.matrix(single_chain_ula$samples)
gradients <- single_chain_ula$der_loglike + single_chain_ula$der_logprior

KSD(samples, gradients)$ksd
```

```
## [1] 2.943607
```

```
imqKSD(samples, gradients)
```

```
## [1] 1.710645
```

```
h_value <- 1
# Running ULA
single_chain_ula <- ULA_fn(d = 11, initial = initial, cov_ULA = cov_ula, h = h_value,
                           iters = its, der_loglike = der_loglike,
                           der_logprior = der_logprior,
                           options = options, varNames = varNames)
```



```

samples <- as.matrix(single_chain_ula$samples)
gradients <- single_chain_ula$der_loglike + single_chain_ula$der_logprior

KSD(samples, gradients)$ksd

## [1] 0.4114754

imqKSD(samples, gradients)

## [1] 0.7582476

h_value <- 1.6
# Running ULA
single_chain_ula <- ULA_fn(d = 11, initial = initial, cov_ULA = cov_ula, h = h_value,
                           iters = its, der_loglike = der_loglike,
                           der_logprior = der_logprior,
                           options = options, varNames = varNames)
samples <- as.matrix(single_chain_ula$samples)
gradients <- single_chain_ula$der_loglike + single_chain_ula$der_logprior

KSD(samples, gradients)$ksd

## [1] 6.967222

imqKSD(samples, gradients)

## [1] 1.599452

```

Investigate convergence for MALA

Try alternative initialisation points you may use the following to simulate from the prior:

```

initial <- recapture_simprior(1, options)

set.seed(2)
n_reps <- 10 # number of chains
its <- 500 # number of MCMC iterations
chains_ULA <- chains_MALA <- samples_ULA <-
  samples_MALA <- list()
for (i in 1:n_reps){
  initial <- recapture_simprior(1, options)
  # Running MALA
  single_chain_mala <- MALA_fn(d = 11, initial = initial, covmala = cov_rw, h = h_mala,
                              iters = its, der_loglike = der_loglike,
                              der_logprior = der_logprior,
                              options = options, varNames = varNames)

  chains_MALA[[i]] <- single_chain_mala
  samples_MALA[[i]] <- single_chain_mala$samples

  # Running ULA
  single_chain_ula <- ULA_fn(d = 11, initial = initial, cov_ULA = cov_ula, h = h_ula,
                             iters = its, der_loglike = der_loglike,
                             der_logprior = der_logprior,
                             options = options, varNames = varNames)

  chains_ULA[[i]] <- single_chain_ula
  samples_ULA[[i]] <- single_chain_ula$samples
}

```

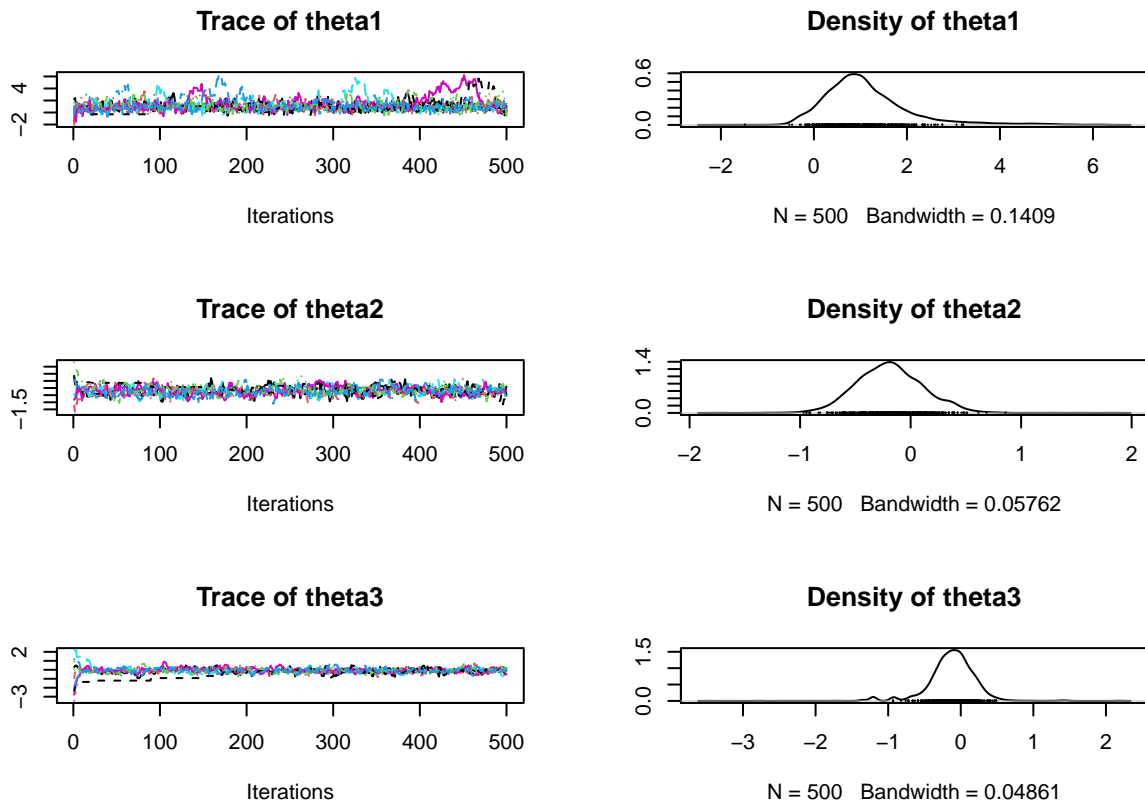
```

}
ula_noburnin <- as.mcmc.list(samples_ULA)
mala_noburnin <- as.mcmc.list(samples_MALA)


```

Investigate the convergence using multiple chains and Gelman & Rubin's \hat{R} diagnostic. Work out a good burnin and investigate control variates on the burnin chain.


```
plot(mala_noburnin, smooth=FALSE)
```




Trace of theta4




Density of theta4




Trace of theta5




Density of theta5




Trace of theta6




Density of theta6




Trace of theta7




Density of theta7




Trace of theta8




Density of theta8







Trace of theta9



Density of theta9





Rhat

The \hat{R} diagnostic is not < 1.1 for all dimensions, so we have evidence of non-convergence.

```
gelman.diag(mala_noburnin, autoburnin=FALSE)
```







```
## Potential scale reduction factors:
##
##      Point est. Upper C.I.
## theta1      1.03      1.04
## theta2      1.02      1.04
## theta3      1.11      1.22
## theta4      1.80      5.48
## theta5      1.70      3.80
## theta6      1.02      1.03
## theta7      1.01      1.02
## theta8      1.01      1.03
## theta9      1.17      1.33
## theta10     1.31      1.62
## theta11     1.12      1.24
##
## Multivariate psrf
##
## 1.57
```

Removing Burn-In


Let's see if we still have evidence of non-convergence if we remove the first 500 samples as burn-in.

```
# Removing burn-in
burnin <- 300
its <- 500
mala_burnin <- list()
for (i in 1:n_reps){
  mala_burnin[[i]] <- mcmc(mala_noburnin[[i]][(burnin+1):(its),])
}
mala_burnin <- as.mcmc.list(mala_burnin)


# Plotting chains
plot(mala_burnin, smooth=FALSE)
```




Trace of theta4




Density of theta4




Trace of theta5




Density of theta5




Trace of theta6




Density of theta6




Trace of theta7




Density of theta7




Trace of theta8




Density of theta8







Trace of theta9



Density of theta9





Calculating Rhat

```
gelman.diag(mala_burnin, autoburnin=FALSE)
```

```
## Potential scale reduction factors:
```

```
##
```

```
##      Point est. Upper C.I.
```

```
## theta1      1.09      1.19
```

```
## theta2      1.03      1.06
```

```
## theta3      1.02      1.05
```

```
## theta4      1.03      1.07
```

```
## theta5      1.04      1.08
```

```
## theta6      1.04      1.07
```

```
## theta7      1.03      1.05
```

```
## theta8      1.02      1.04
```

```
## theta9      1.03      1.06
```

```
## theta10     1.02      1.05
```

```
## theta11     1.01      1.03
```

```
##
```

```
## Multivariate psrf
```

```
##
```

```
## 1.09
```

Investigate the KSD without bias:

```
burnin <- 300
```

```
KSD_MALA_gaussian <- KSD_ULA_gaussian <- rep(NA,n_reps)
```

```
KSD_MALA_imq <- KSD_ULA_imq <- rep(NA,n_reps)
```

```


for (i in 1:n_reps){
  samples <- as.matrix(chains_MALA[[i]]$samples[(burnin+1):(its),])
  gradients <- chains_MALA[[i]]$der_loglike[(burnin+1):(its),] + chains_MALA[[i]]$der_logprior[(burnin+1):(its),]

  KSD_MALA_gaussian[i] <- KSD(samples, gradients)$ksd
  KSD_MALA_imq[i] <- imqKSD(samples, gradients)

  samples <- as.matrix(chains_ULA[[i]]$samples[(burnin+1):(its),])
  gradients <- chains_ULA[[i]]$der_loglike[(burnin+1):(its),] + chains_ULA[[i]]$der_logprior[(burnin+1):(its),]

  KSD_ULA_gaussian[i] <- KSD(samples, gradients)$ksd
  KSD_ULA_imq[i] <- imqKSD(samples, gradients)
}
boxplot(KSD_MALA_gaussian, KSD_ULA_gaussian)

```



```

boxplot(KSD_MALA_imq, KSD_ULA_imq)

```

